


CS 188: Artificial Intelligence Fall 2010

Lecture 20: Particle Filtering 11/4/2010


Dan Klein – UC Berkeley

Recap: Reasoning Over Time

Markov models



Hidden Markov models




Recap: Filtering

Elapse time: compute $P(X_t | e_{1:t-1})$

$$P(x_t | e_{1:t-1}) = \sum_{x_{t-1}} P(x_{t-1} | e_{1:t-1}) \cdot P(x_t | x_{t-1})$$

Observe: compute $P(X_t | e_{1:t})$

$$P(x_t | e_{1:t}) \propto P(x_t | e_{1:t-1}) \cdot P(e_t | x_t)$$



Particle Filtering


Sometimes $|X|$ is too big to use exact inference

- $|X|$ may be too big to even store $B(X)$
- E.g. X is continuous
- $|X|^2$ may be too big to do updates

0.0	0.1	0.0
0.0	0.0	0.2
0.0	0.2	0.5

Solution: approximate inference

- Track samples of X , not all values
- Samples are called particles
- Time per step is linear in the number of samples
- But: number needed may be large
- In memory: list of particles, not states



This is how robot localization works in practice

Representation: Particles


- Our representation of $P(X)$ is now a list of N particles (samples)

- Generally, $N \ll |X|$
- Storing map from X to counts would defeat the point

- $P(x)$ approximated by number of particles with value x

- So, many x will have $P(x) = 0$!
- More particles, more accuracy

- For now, all particles have a weight of 1



Particles:
(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(2,1)
(3,3)
(3,3)
(2,1)

6

Particle Filtering: Elapse Time


- Each particle is moved by sampling its next position from the transition model

$$x' = \text{sample}(P(X'|x))$$

- This is like prior sampling – samples' frequencies reflect the transition probs
- Here, most samples move clockwise, but some move in another direction or stay in place

- This captures the passage of time

- If we have enough samples, close to the exact values before and after (consistent)



Particle Filtering: Observe


- Slightly trickier:

- Don't do rejection sampling (why not?)
- We don't sample the observation, we fix it
- This is similar to likelihood weighting, so we downweight our samples based on the evidence

$$w(x) = P(e|x)$$

$$B(X) \propto P(e|X)B'(X)$$

- Note that, as before, the probabilities don't sum to one, since most have been downweighted (in fact they sum to an approximation of $P(e)$)



Particle Filtering: Resample

- Rather than tracking weighted samples, we resample

- N times, we choose from our weighted sample distribution (i.e. draw with replacement)

- This is analogous to renormalizing the distribution


- Now the update is complete for this time step, continue with the next one

Old Particles:

(3,3) $w=0.1$
(2,1) $w=0.9$
(2,1) $w=0.9$
(3,1) $w=0.4$
(3,2) $w=0.3$
(2,2) $w=0.4$
(1,1) $w=0.4$
(3,1) $w=0.4$
(2,1) $w=0.9$
(3,2) $w=0.3$

New Particles:

(2,1) $w=1$
(2,1) $w=1$
(2,1) $w=1$
(3,2) $w=1$
(2,2) $w=1$
(2,1) $w=1$
(1,1) $w=1$
(3,1) $w=1$
(2,1) $w=1$
(1,1) $w=1$




Robot Localization

- In robot localization:

- We know the map, but not the robot's position
- Observations may be vectors of range finder readings
- State space and readings are typically continuous (works basically like a very fine grid) and so we cannot store $B(X)$
- Particle filtering is a main technique

- [Demos]



P4: Ghostbusters 2.0 (beta)


- Plot:** Pacman's grandfather, Grandpac, learned to hunt ghosts for sport.

- He was blinded by his power, but could hear the ghosts' banging and clanging.

- Transition Model:** All ghosts move randomly, but are sometimes biased

- Emission Model:** Pacman knows a "noisy" distance to each ghost


Noisy distance prob
True distance = 8



[Demo]

Dynamic Bayes Nets (DBNs)


- We want to track multiple variables over time, using multiple sources of evidence
- Idea: Repeat a fixed Bayes net structure at each time
- Variables from time t can condition on those from $t-1$



- DBNs with evidence at leaves are (in principle) HMMs

Exact Inference in DBNs

- Variable elimination applies to dynamic Bayes nets
- Procedure: "unroll" the network for T time steps, then eliminate variables until $P(X_T | e_{1:T})$ is computed



- Online belief updates: Eliminate all variables from the previous time step; store factors for current time only

13

DBN Particle Filters

- A particle is a **complete** sample for a time step
- Initialize**: Generate prior samples for the $t=1$ Bayes net
 - Example particle: $\mathbf{G}_1^a = (3,3)$ $\mathbf{G}_1^b = (5,3)$
- Elastpse time**: Sample a successor for each particle
 - Example successor: $\mathbf{G}_2^a = (2,3)$ $\mathbf{G}_2^b = (6,3)$
- Observe**: Weight each entire sample by the likelihood of the evidence conditioned on the sample
 - Likelihood: $P(\mathbf{E}_1^a | \mathbf{G}_1^a) * P(\mathbf{E}_1^b | \mathbf{G}_1^b)$
- Resample**: Select prior samples (tuples of values) in proportion to their likelihood

[Demo]

SLAM

- SLAM = Simultaneous Localization And Mapping**
 - We do not know the map or our location
 - Our belief state is over maps and positions!
 - Main techniques: Kalman filtering (Gaussian HMMs) and particle methods




DP-SLAM, Ron Parr

[DEMOS]

HMMs: MLE Queries

- HMMs defined by

- States X
- Observations E
- Initial distr: $P(X_1)$
- Transitions: $P(X|X_{-1})$
- Emissions: $P(E|X)$




- Query: most likely explanation:

$$\arg \max_{x_{1:t}} P(x_{1:t} | e_{1:t})$$

16

State Path Trellis


- State trellis: graph of states and transitions over time



- Each arc represents some transition $x_{t-1} \rightarrow x_t$
- Each arc has weight $P(x_t | x_{t-1})P(e_t | x_t)$
- Each path is a sequence of states
- The product of weights on a path is the seq's probability
- Can think of the Forward (and now Viterbi) algorithms as computing sums of all paths (best paths) in this graph

17

Viterbi Algorithm



$$x_{1:T}^* = \arg \max_{x_{1:T}} P(x_{1:T} | e_{1:T}) = \arg \max_{x_{1:T}} P(x_{1:T}, e_{1:T})$$

$$m_t[x_t] = \max_{x_{1:t-1}} P(x_{1:t-1}, x_t, e_{1:t})$$


$$= \max_{x_{1:t-1}} P(x_{1:t-1}, e_{1:t-1}) P(x_t | x_{t-1}) P(e_t | x_t)$$

$$= P(e_t | x_t) \max_{x_{t-1}} P(x_t | x_{t-1}) \max_{x_{1:t-2}} P(x_{1:t-1}, e_{1:t-1})$$

$$= P(e_t | x_t) \max_{x_{t-1}} P(x_t | x_{t-1}) m_{t-1}[x_{t-1}]$$

18

Example



19